

POSIX vs. Win32 Threads

In depth threading techniques

Disclaimer

- ▶ This presentation provides a summary based on few articles regarding threads in Win32 and POSIX systems. All these articles are mentioned in the ***References*** slide.

What is Thread

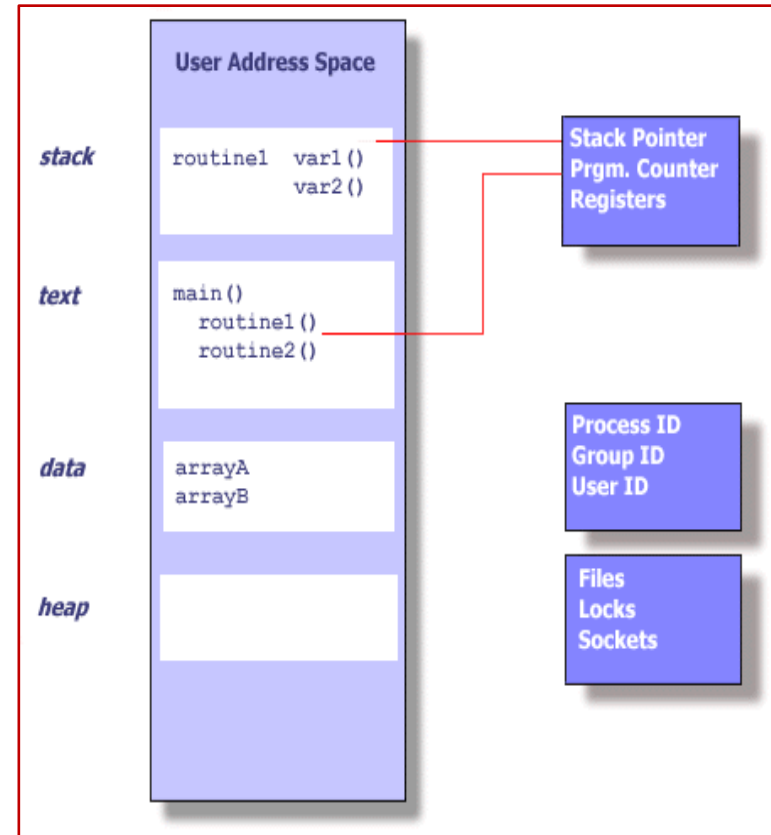
1. Processes

▶ What is a Process

- A process is created **by OS** to run your **Main** program.

▶ Contains all information to be handled by the OS.

- Process ID, process group ID, user ID, and group ID
- Environment
- **Scheduling Properties (priority, etc.)**
- Working directory.
- Program instructions
- **Registers, Stack, Heap**
- File descriptors
- **Signal actions**
- Shared libraries
- Inter-process communication tools
shared memory , message queues,
semaphores , pipes,



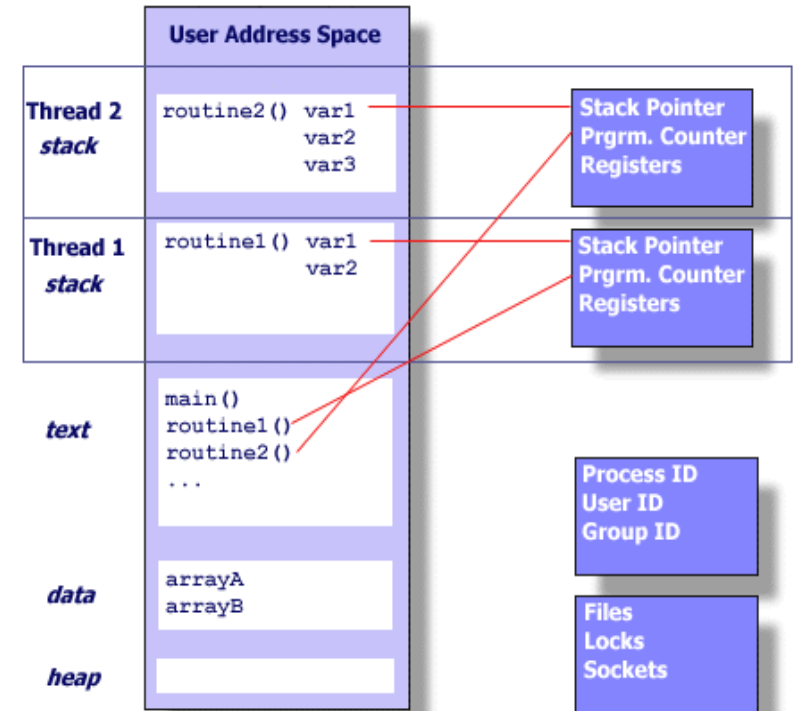
What is Thread

2. Threads

- What is a Thread
 - A thread is created **by OS** to run a stream of instructions.
- Threads contain smaller set of information
 - *Scheduling Properties*
 - *Registers*
 - *Stack*
 - *Signal actions*
- Light-Weight

Most overhead accomplished through maintaining its process.

 - Exists within a process and uses the process resources
 - Has its own independent flow of control (as long as its parent process exists and the OS supports it)
 - Duplicates only essential resources for independent scheduling



Comparing timing results in POSIX for the **fork()** and **pthread_create()** (reflecting 50,000 creations) shows **30-60 times overhead!**

Threads Sharing

➤ **Resource sharing.**

Changes made by one thread to shared system resources will be seen by all other threads.

- File open or close at the process level.

➤ **Memory Sharing (Same address, same data).**

Two pointers having the same value, exist in the same virtual memory and point to the same data.

- No copy is needed!

➤ **Synchronization required.**

As resources and memory sharing is possible, explicit synchronization is required by the programmer

Win32 vs. POSIX Interface

1. Function calls

Win32 Threads	<u>PThreads</u>
Just one type: HANDLE	Each object has its own data type : <u>pthread_t</u> , <u>pthread_mutex_t</u> , ...
One function needed to make one functionality (e.g. <u>WaitForSingleObject</u>)	Each object has its own functions/attributes
Simpler and more generic - is it OOP? (sounds like misusing void* in C).	Reading and understanding may be more straightforward and less confusing.

Win32 vs. POSIX Interface

2. Synchronization overhead

Win32 Threads	<u>PThreads</u>
Synchronization objects: <ul style="list-style-type: none">– Events– Semaphores– Mutexes– Critical sections	Synchronization primitives: <ul style="list-style-type: none">– Semaphores– Mutexes– Conditional variables
<p><i>once an event is in the signaled state, it stays signaled.</i></p> <p>However, It is up to the programmer to ensure the proper switching of Windows events from the signaled to <u>unsignaled</u> state. (When an object is signaled, you have to check what other objects in the awoken thread might be waiting for and remove it from those wait queues)</p>	<p><i>Signals to condition variables are either "caught" by waiting thread(s) or discarded.</i></p> <p>However, use of a well known coding structure at each access of a condition variable will ensure no signals are "lost" by threads that may not be waiting at the exact time of signaling</p>

More issues to consider

- ▶ Historically, hardware vendors have implemented their own proprietary versions of threads.
- ▶ In Windows, the thread is the basic execution unit, and the process is a container that holds this thread.
In Linux, the basic execution unit is the process
- ▶ Some may claim POSIX threads are a low-level API and Windows threads are a high-level API
- ▶ In Windows the thread scheduling code is implemented in the kernel.
- ▶ In Linux realm, there are several drafts of the POSIX threads standard. Differences between drafts exist! Especially many scheduling mechanisms exist (user and kernel-level threads).
- ▶ The current POSIX standard is defined only for the C language.

Mapping WIN32 to PTHREADS

1. Process Mapping

Win32	Linux	Classification
CreateProcess() CreateProcessAsUser()	fork() setuid() exec()	Mappable
TerminateProcess()	kill()	Mappable
SetThreadpriority() GetThreadPriority()	Setpriority() getPriority()	Mappable
GetCurrentProcessID()	getpid()	Mappable
Exitprocess()	exit()	Mappable
Waitforsingleobject() Waitformultipleobject() GetExitCodeProcess()	waitpid() !! Using semaphores Waitforsingleobject / Waitformultipleobject can be implemented	Context specific
GetEnvironmentVariable SetEnvironmentVariable	getenv() setenv()	Mappable

Mappable: Both the Windows and Linux constructs provide similar functionality.

Context: The decision to use a specific Linux construct(s) depends on the application context

**!! More on WaitForMultipleObjects may be found under the following
reference:** <http://www.ibm.com/developerworks/linux/library/l-ipc2lin3.html>

Mapping WIN32 to PTHREADS

2. Thread Mapping

Win32	Linux	Classification
CreateThread	pthread_create pthread_attr_init pthread_attr_setstacksize pthread_attr_destroy	Mappable
ThreadExit	pthread_exit	Mappable
WaitForSingleObject	pthread_join pthread_attr_setdetachstate pthread_detach	Mappable
SetPriorityClass SetThreadPriority	setpriority sched_setscheduler sched_setparam pthread_setschedparam pthread_setschedpolicy pthread_attr_setschedparam pthread_attr_setschedpolicy	Context Specific

Mappable: Both the Windows and Linux constructs provide similar functionality.

Context: The decision to use a specific Linux construct(s) depends on the application context

Mapping WIN32 to PTHREADS

3. Synchronization

Win32 Thread Level	Linux Thread Level	Linux Process Level
Mutex	Mutex - pthread library	System V semaphores
Critical section	Mutex - pthread library	N/A <i>Critical sections are used only within threads of the same process</i>
Semaphore	Conditional Variable with mutex – pthreads POSIX semaphores	System V Semaphores
Event	Conditional Variable with mutex – pthreads POSIX semaphores	System V Semaphores

!! **Conditional Variables:** While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.

See more under: <https://computing.llnl.gov/tutorials/pthreads>

Mapping WIN32 to PTHREADS

Synchronization (semaphores)

Win32 Thread	Linux Thread	Linux Process	Classification
CreateSemaphore	sem_init	semget semctl	Context specific
OpenSemaphore	N/A	semget	Context specific
WaitForSingleObject	sem_wait sem_trywait	semop	Context specific
ReleaseSemaphore	sem_post	semop	Context specific
CloseHandle	sem_destroy	semctl	Context specific



Notice: Win32 semaphores are created within a thread, and propagated all over the system. POSIX semaphores are wither for inter-thread within a process, or inter-process entities.

!! **opensemaphore** function enables multiple processes to open handles of the same semaphore object. The function succeeds only if some process has already created the semaphore by using the **CreateSemaphore** function.

Mapping WIN32 to PTHREADS

Synchronization (events)

Win32 Thread	Linux Thread	Linux Process	Classification
CreateEvent OpenEvent	pthread_cond_init sem_init	semget semctl	context specific
SetEvent	pthread_cond_signal sem_post	semop	context specific
ResetEvent	N/A	N/A	context specific
WaitForSingleObject	pthread_cond_wait pthread_cond_timedwait sem_wait sem_trywait	semop	context specific
CloseHandle	pthread_cond_destroy sem_destroy	semctl	context specific



Notice (See conclusion in the next slide).

- Events are Windows specific objects.
- POSIX semaphores with count set to 1 provide similar functionality to the Windows events. However, they don't provide timeout in the wait functions.
- Conditional variables & Mutex in pthreads provide event-based synchronization between threads, but they are **synchronous**.

Mapping WIN32 to PTHREADS

Synchronization (events) - conclusion



Win32 Threads	PThreads
Both Manual/ Auto-reset events	Only Auto-reset mechanism.
<i>Named</i> events to synchronization processes. Un-named events to synchronize threads.	Synchronization is inter-thread based. (System V semaphore or signals can be used)
Event objects initial state is set to signaled .	Does not provide an initial state,. (<i>POSIX semaphores provide initial state</i>)
Event objects are asynchronous	Conditional variables are synchronous (<i>POSIX/System V events are asynchronous</i>)
Timeout value can be specified	Timeout value can be specified (<i>Other Linux systems don't support timeout</i>)

Mapping WIN32 to PTHREADS

Synchronization (mutex)

Win32 Thread	Linux Thread	Linux Process	Classification
CreateMutex	pthread_mutex_init	semget semctl	context specific
OpenMutex	N/A	semget	context specific
WaitForSingleObject	pthread_mutex_lock pthread_mutex_trylock	semop	context specific
ReleaseMutex	pthread_mutex_unlock	semop	context specific
CloseHandle	pthread_mutex_destroy	semctl	context specific



Notice (See conclusion in the next slide).

Some major differences reside, including:

- Named and un-named mutexes.
- Ownership at creation.
- Timeout during wait.
- Recursive mutexes.

Mapping WIN32 to LINUX

Synchronization (mutex) - conclusion

Win32 Threads	PThreads
<i>Named</i> mutexes synchronize process/thread Un-named mutex synchronize threads.	Synchronization is inter-thread based. (System V semaphore or signals can be used)
Can be owned during creation.	To achieve the same in Linux, a mutex should be locked explicitly after creation.
Recursive by default	<i>Pthread has recursive mutex (initialized explicitly)</i> <i>(other Linux system do not allow recursive mutex)</i>
Timeout value can be specified	Timeout not available. <i>(can be obtained via application logic)</i>

Mapping WIN32 to LINUX

Synchronization (Critical Sections)

Win32 Thread	Linux Thread	Classification
InitializeCriticalSection InitializeCriticalSectionAndSpinCount	pthread_mutex_init	Mappable
EnterCriticalSection TryEnterCriticalSection	pthread_mutex_lock pthread_mutex_trylock	Mappable
LeaveCriticalSection	pthread_mutex_trylock	Mappable
DeleteCriticalSection	pthread_mutex_destroy	Mappable

Notice: Since the critical sections are used only between the threads of the same process, Pthreads mutex can be used to achieve the same **functionality** on Linux systems.



1 - In general, mutex objects are more CPU intensive and slower to use than critical sections due to a larger amount of bookkeeping, and the deep execution path into the kernel taken to acquire a mutex. The equivalent functions for accessing a critical section remain in thread context, and consist of merely checking and incrementing/decrementing lock-count related data. This makes critical sections light weight, fast, and easy to use for thread synchronization within a process.

2- In Win32 The primary benefit of using a mutex object is that a mutex can be named and shared across process boundaries.

References

- ▶ **POSIX Threads Programming**

<https://computing.llnl.gov/tutorials/pthreads/>

- ▶ **Why pthreads are better than Win32**

<http://softwarecommunity.intel.com/ISN/Community/en-US/forums/post/840096.aspx>

Why Windows threads are better than pthreads:

<http://softwareblogs.intel.com/2006/10/19/why-windows-threads-are-better-than-posix-threads/>

- ▶ **Port Windows IPC apps to Linux (including code examples)**

<http://www.ibm.com/developerworks/linux/library/l-ipc2lin1.html>

<http://www.ibm.com/developerworks/linux/library/l-ipc2lin2.html>

<http://www.ibm.com/developerworks/linux/library/l-ipc2lin3.html>

Thank you